

CUDA Atomics

Marc-Antoine Le Guen



Atomics

- Algunas situaciones son muy simple en single-thread
 - pero pueden representar un desafío muy grande implementarlo en un programa masivamente paralelo.



Atomics - Objetivos

- Entender que son las atomics operations y para qué son útiles
- Operaciones aritméticas con operaciones atómicas



Atomics operations

- read-modify-write operation
 - `x++`
 - Quisiéramos :

Table 9.2 Two threads incrementing the value in `x`

STEP	EXAMPLE
1. Thread A reads the value in <code>x</code> .	A reads 7 from <code>x</code> .
2. Thread A adds 1 to the value it read.	A computes 8.
3. Thread A writes the result back to <code>x</code> .	<code>x <- 8.</code>
4. Thread B reads the value in <code>x</code> .	B reads 8 from <code>x</code> .
5. Thread B adds 1 to the value it read.	B computes 9.
6. Thread B writes the result back to <code>x</code> .	<code>x <- 9.</code>

Atomics operations

- read-modify-write operation
 - `x++`
 - Podemos obtener :

Table 9.3 Two threads incrementing the value in `x` with interleaved operations

STEP	EXAMPLE
Thread A reads the value in <code>x</code> .	A reads 7 from <code>x</code> .
Thread B reads the value in <code>x</code> .	B reads 7 from <code>x</code> .
Thread A adds 1 to the value it read.	A computes 8.
Thread B adds 1 to the value it read.	B computes 8.
Thread A writes the result back to <code>x</code> .	<code>x <- 8.</code>
Thread B writes the result back to <code>x</code> .	<code>x <- 8.</code>

Atomics operations

- Las operaciones que no descomponen el proceso de read-modify-write son **atomics** **operaciones**.



Atomics - Ejemplo : Histograma

- Contar la frecuencia de cada dato en una serie de datos
 - colores de píxeles en una imagen
 - letras en un texto

2	2	1	2	1	2	2	1	1	1	2	1	1	1
A	C	D	G	H	I	M	N	O	P	R	T	U	W

Figure 9.1 Letter frequency histogram built from the string *Programming with CUDA C*



Atomics - Ejemplo : Histograma

- Muy simple en CPU
 - Inicialización de datos 100 MB
 - 256 - 8 bits
 - histo 256 inicializado a 0

```
#define SIZE  (100*1024*1024)

int main( void ) {
    unsigned char *buffer = (unsigned char*)big_random_block( SIZE );

    unsigned int  histo[256];
    for (int i=0; i<256; i++)
        histo[i] = 0;
```



Atomics - Ejemplo : Histograma

- Muy simple en CPU
 - generar el histograma
 - 0.31 segundos CPU

```
for (int i=0; i<SIZE; i++)  
    histo[buffer[i]]++;
```

- Suma del histograma = cantidad de datos

```
long histoCount = 0;  
for (int i=0; i<256; i++) {  
    histoCount += histo[i];  
}
```



Atomics - Ejemplo : Histograma

- GPU
- Inicialización clásica
 - `cudaMemset`

```
int main( void ) {  
    unsigned char *buffer =(unsigned char*)big_random_block( SIZE );  
  
    // allocate memory on the GPU for the file's data  
    unsigned char *dev_buffer;  
    unsigned int *dev_histo;  
    cudaMalloc( (void**)&dev_buffer, SIZE );  
    cudaMemcpy( dev_buffer, buffer, SIZE, cudaMemcpyHostToDevice );  
  
    cudaMalloc( (void**)&dev_histo, 256 * sizeof( int ) );  
    cudaMemset( dev_histo, 0, 256 * sizeof( int ) );  
}
```

Atomics - Ejemplo : Histograma

- GPU
- Inicialización clásica
 - `cudaMemset`

```
int main( void ) {  
    unsigned char *buffer =(unsigned char*)big_random_block( SIZE );  
  
    // allocate memory on the GPU for the file's data  
    unsigned char *dev_buffer;  
    unsigned int *dev_histo;  
    cudaMalloc( (void**)&dev_buffer, SIZE );  
    cudaMemcpy( dev_buffer, buffer, SIZE, cudaMemcpyHostToDevice );  
    atomicAdd
```



Atomics - Ejemplo : Histograma

- GPU
- Inicialización clásica
 - `cudaMemset`

```
int main( void ) {  
    unsigned char *buffer =(unsigned char*)big_random_block( SIZE );  
  
    // allocate memory on the GPU for the file's data  
    unsigned char *dev_buffer;  
    unsigned int *dev_histo;  
    cudaMalloc( (void**)&dev_buffer, SIZE );  
    cudaMemcpy( dev_buffer, buffer, SIZE, cudaMemcpyHostToDevice );  
  
    cudaMalloc( (void**)&dev_histo, 256 * sizeof( int ) );  
    cudaMemset( dev_histo, 0, 256 * sizeof( int ) );  
}
```



Atomics - Ejemplo : Histograma

- GPU
- Inicialización clásica
 - Asignación de memoria en el Host para después recuperar el histograma resultado

```
unsigned int histo[256];  
...  
cudaMemcpy( histo, dev_histo, 256 * sizeof( int ),  
cudaMemcpyDeviceToHost );
```



Atomics - Ejemplo : Histograma

- GPU
- **Verificación de los resultados**
 - Calcular en el CPU un histograma al reverse
 - restar / sumar
 - Si existe una elemento diferente de 0 en el histograma, nuestro programa GPGPU fallo.

```
for (int i=0; i<SIZE; i++)  
    histo[buffer[i]]--;  
for (int i=0; i<256; i++) {  
    if (histo[i] != 0)  
        printf( "Failure at %d!  Off by %d\n", i, histo[i] );  
}
```



Atomics - Ejemplo : Histograma

- Kernel

- `atomicAdd(addr , value)`
 - el gpu nos asegura la conservación de la unidad del proceso read-modify-write

```
__global__ void histo_kernel( unsigned char *buffer,
                             long size,
                             unsigned int *histo ) {
    // calculate the starting index and the offset to
    the next
    // block that each thread will be processing
    int i = threadIdx.x + blockIdx.x * blockDim.x;
    int stride = blockDim.x * gridDim.x;
    while (i < size) {
        atomicAdd( &histo[buffer[i]], 1 );
        i += stride;
    }
}
```


Atomics - Ejemplo : Histograma

- Problema, este algoritmo produce un problema debido a la gran cantidad de threads en competición para acceder a la misma memoria
 - Muchas operaciones serializadas.
 - Verdad en GPU de compute capability ~ 1.3
 - 0.31 segundos CPU
 - 1.7 segundos 285 gtx 1.3 vs 0.035 segundos 670 gtx 3.0

```
__global__ void histo_kernel( unsigned char *buffer,
                             long size,
                             unsigned int *histo ) {
    // calculate the starting index and the offset to the next
    // block that each thread will be processing
    int i = threadIdx.x + blockIdx.x * blockDim.x;
    int stride = blockDim.x * gridDim.x;
    while (i < size) {
        atomicAdd( &histo[buffer[i]], 1 );
        i += stride;
    }
}
```

Atomics - Ejemplo : Histograma

```
int main( void ) {
    unsigned char *buffer =
        (unsigned char*)big_random_block ( SIZE );

    // capture the start time
    // starting the timer here so that we include the cost of
    // all of the operations on the GPU.
    cudaEvent_t start, stop;
    cudaEventCreate ( &start );
    cudaEventCreate ( &stop );
    cudaEventRecord ( start, 0 );
    // allocate memory on the GPU for the file's data
    unsigned char *dev_buffer;
    unsigned int *dev_histo;
    cudaMalloc ( (void**)&dev_buffer, SIZE );
    cudaMemcpy ( dev_buffer, buffer, SIZE,
        cudaMemcpyHostToDevice );

    cudaMalloc ( (void**)&dev_histo,
        256 * sizeof( int ) );
    cudaMemset ( dev_histo, 0,
        256 * sizeof( int ) );

    // kernel launch - 2x the number of mps gave best timing
    cudaDeviceProp prop;
    cudaGetDeviceProperties ( &prop, 0 );
    int blocks = prop.multiProcessorCount;
    printf(" multiprocesseur : %d", blocks);
    histo_kernel <<<blocks*2,256>>>( dev_buffer, SIZE, dev_histo
);

    unsigned int histo[256];
    cudaMemcpy ( histo, dev_histo,
        256 * sizeof( int ),
        cudaMemcpyDeviceToHost );
}
```

```
// get stop time, and display the timing results
    cudaEventRecord ( stop, 0 );
    cudaEventSynchronize ( stop );
    float elapsedTime;
    cudaEventElapsedTime ( &elapsedTime,
        start, stop );

    printf ( "Time to generate: %3.1f ms\n", elapsedTime );

    long histoCount = 0;
    for (int i=0; i<256; i++) {
        histoCount += histo[i];
    }
    printf ( "Histogram Sum: %ld\n", histoCount );

    // verify that we have the same counts via CPU
    for (int i=0; i<SIZE; i++)
        histo[buffer[i]]--;
    for (int i=0; i<256; i++) {
        if (histo[i] != 0)
            printf ( "Failure at %d! Off by %d\n", i, histo[i] );
    }

    cudaEventDestroy ( start );
    cudaEventDestroy ( stop );
    cudaFree ( dev_histo );
    cudaFree ( dev_buffer );
    free( buffer );
    system("PAUSE");
    return 0;
}
```

Atomics - Ejemplo : Histograma

```
int main( void ) {
    unsigned char *buffer =
        (unsigned char*)big_random_block ( SIZE );

    // capture the start time
    // starting the timer here so that we include the cost of
    // all of the operations on the GPU.
    cudaEvent_t start, stop;
    cudaEventCreate ( &start );
    cudaEventCreate ( &stop );
    cudaEventRecord ( start, 0 );
    // allocate memory on the GPU for the file's data
    unsigned char *dev_buffer;
    unsigned int *dev_histo;
    cudaMalloc ( (void**)&dev_buffer, SIZE );
    cudaMemcpy ( dev_buffer, buffer, SIZE,
                cudaMemcpyHostToDevice );

    cudaMalloc ( (void**)&dev_histo,
                256 * sizeof( int ) );
    cudaMemset ( dev_histo, 0,
                256 * sizeof( int ) );

    // kernel launch - 2x the number of mps gave best timing
    cudaDeviceProp prop;
    cudaGetDeviceProperties ( &prop, 0 );
    int blocks = prop.multiProcessorCount;
    printf(" multiprocesseur : %d" , blocks);
    histo_kernel <<<blocks*2,256>>>( dev_buffer, SIZE, dev_histo
);

    unsigned int histo[256];
    cudaMemcpy ( histo, dev_histo,
                256 * sizeof( int ),
                cudaMemcpyDeviceToHost );
}
```

```
// get stop time, and display the timing results
    cudaEventRecord ( stop, 0 );
    cudaEventSynchronize ( stop );
    float elapsedTime;
    cudaEventElapsedTime ( &elapsedTime,
                          start, stop );
    printf( "Time to generate: %3.1f ms\n" , elapsedTime );

    long histoCount = 0;
    for (int i=0; i<256; i++) {
        histoCount += histo[i];
    }
    printf( "Histogram Sum: %ld\n" , histoCount );

    // verify that we have the same counts via CPU
    for (int i=0; i<SIZE; i++)
        histo[buffer[i]]--;
    for (int i=0; i<256; i++) {
        if (histo[i] != 0)
            printf ( "Failure at %d! Off by %d\n" , i, histo[i] );
    }

    cudaEventDestroy ( start );
    cudaEventDestroy ( stop );
    cudaFree ( dev_histo );
    cudaFree ( dev_buffer );
    free( buffer );
    system("PAUSE");
    return 0;
}
```

Atomics - Ejemplo : Histograma

```
int main( void ) {
    unsigned char *buffer =
        (unsigned char*)big_random_block ( SIZE );

    // capture the start time
    // starting the timer here so that we include the cost of
    // all of the operations on the GPU.
    cudaEvent_t start, stop;
    cudaEventCreate ( &start );
    cudaEventCreate ( &stop );
    cudaEventRecord ( start, 0 );
    // allocate memory on the GPU for the file's data
    unsigned char *dev_buffer;
    unsigned int *dev_histo;
    cudaMalloc ( (void**)&dev_buffer, SIZE );
    cudaMemcpy ( dev_buffer, buffer, SIZE,
                cudaMemcpyHostToDevice );

    cudaMalloc ( (void**)&dev_histo,
                256 * sizeof( int ) );
    cudaMemset ( dev_histo, 0,
                256 * sizeof( int ) );

    // kernel launch - 2x the number of mps gave best timing
    cudaDeviceProp prop;
    cudaGetDeviceProperties ( &prop, 0 );
    int blocks = prop.multiProcessorCount;
    histo_kernel<<<blocks*2,256>>>( dev_buffer, SIZE, dev_histo );

    unsigned int histo[256];
    cudaMemcpy ( histo, dev_histo,
                256 * sizeof( int ),
                cudaMemcpyDeviceToHost );
}
```

Atomics - Ejemplo : Histograma

```
int main( void ) {
    unsigned char *buffer =
        (unsigned char*)big_random_block ( SIZE );

    // capture the start time
    // starting the timer here so that we include the cost of
    // all of the operations on the GPU.
    cudaEvent_t start, stop;
    cudaEventCreate ( &start );
    cudaEventCreate ( &stop );
    cudaEventRecord ( start, 0 );
    // allocate memory on the GPU for the file's data
    unsigned char *dev_buffer;
    unsigned int *dev_histo;
    cudaMalloc ( (void**)&dev_buffer, SIZE );
    cudaMemcpy ( dev_buffer, buffer, SIZE,
        cudaMemcpyHostToDevice );

    cudaMalloc ( (void**)&dev_histo,
        256 * sizeof( int ) );
    cudaMemset ( dev_histo, 0,
        256 * sizeof( int ) );

    // kernel launch - 2x the number of mps gave best timing
    cudaDeviceProp prop;
    cudaGetDeviceProperties ( &prop, 0 );
    int blocks = prop.multiProcessorCount;
    printf(" multiprocesseur : %d", blocks);
    histo_kernel <<<blocks*2,256>>>( dev_buffer, SIZE, dev_histo
);

    unsigned int histo[256];
    cudaMemcpy ( histo, dev_histo,
        256 * sizeof( int ),
        cudaMemcpyDeviceToHost );
```

```
// get stop time, and display the timing results
    cudaEventRecord ( stop, 0 );
    cudaEventSynchronize ( stop );
    float elapsedTime;
    cudaEventElapsedTime ( &elapsedTime,
        start, stop );
    printf ( "Time to generate: %3.1f ms\n" , elapsedTime );

    long histoCount = 0;
    for (int i=0; i<256; i++) {
        histoCount += histo[i];
    }
    printf ( "Histogram Sum: %ld\n" , histoCount );

    // verify that we have the same counts via CPU
    for (int i=0; i<SIZE; i++)
        histo[buffer[i]]--;
    for (int i=0; i<256; i++) {
        if (histo[i] != 0)
            printf ( "Failure at %d! Off by %d\n" , i, histo[i] );
    }

    cudaEventDestroy ( start );
    cudaEventDestroy ( stop );
    cudaFree ( dev_histo );
    cudaFree ( dev_buffer );
    free( buffer );
    system("PAUSE");
    return 0;
}
```

Atomics - Ejemplo : Histograma

```
int main( void ) {
    unsigned char *buffer =
        (unsigned char*)big_random_block ( SIZE );

    // capture the start time
    // starting the timer here so that we include the cost of
    // all of the operations on the GPU.
    cudaEvent_t start, stop;
    cudaEventCreate ( &start );
    cudaEventCreate ( &stop );
    cudaEventRecord ( start, 0 );
    // allocate memory on the GPU for the file's data
    unsigned char *dev_buffer;
    unsigned int *dev_histo;
    cudaMalloc ( (void**)&dev_buffer, SIZE );
    cudaMemcpy ( dev_buffer, buffer, SIZE,
        cudaMemcpyHostToDevice );

    cudaMalloc ( (void**)&dev_histo,
        256 * sizeof( int ) );
    cudaMemset ( dev_histo, 0,
        256 * sizeof( int ) );

    // kernel launch - 2x the number of mps gave best timing
    cudaDeviceProp prop;
    cudaGetDeviceProperties ( &prop, 0 );
    int blocks = prop.multiProcessorCount;
    printf(" multiprocesseur : %d", blocks);
    histo_kernel<<<blocks*2,256>>>( dev_buffer, SIZE, dev_histo
);

    unsigned int histo[256];
    cudaMemcpy ( histo, dev_histo,
        256 * sizeof( int ),
        cudaMemcpyDeviceToHost );
};
```

```
// get stop time, and display the timing results
    cudaEventRecord ( stop, 0 );
    cudaEventSynchronize ( stop );
    float elapsedTime;
    cudaEventElapsedTime ( &elapsedTime,
        start, stop );
    printf( "Time to generate: %3.1f ms\n", elapsedTime );

    long histoCount = 0;
    for (int i=0; i<256; i++) {
        histoCount += histo[i];
    }
    printf( "Histogram Sum: %ld\n", histoCount );

    // verify that we have the same counts via CPU
    for (int i=0; i<SIZE; i++)
        histo[buffer[i]]--;
    for (int i=0; i<256; i++) {
        if (histo[i] != 0)
            printf ( "Failure at %d! Off by %d\n", i, histo[i] );
    }

    cudaEventDestroy ( start );
    cudaEventDestroy ( stop );
    cudaFree ( dev_histo );
    cudaFree ( dev_buffer );
    free( buffer );
    system("PAUSE");
    return 0;
};
```

Atomics - Ejemplo : Histograma

```
int main( void ) {
    unsigned char *buffer =
        (unsigned char*)big_random_block ( SIZE );

    // capture the start time
    // starting the timer here so that we include the cost of
    // all of the operations on the GPU.
    cudaEvent_t start, stop;
    cudaEventCreate ( &start );
    cudaEventCreate ( &stop );
    cudaEventRecord ( start, 0 );
    // allocate memory on the GPU for the file's data
    unsigned char *dev_buffer;
    unsigned int *dev_histo;
    cudaMalloc ( (void**)&dev_buffer, SIZE );
    cudaMemcpy ( dev_buffer, buffer, SIZE,
        cudaMemcpyHostToDevice );

    cudaMalloc ( (void**)&dev_histo,
        256 * sizeof( int ) );
    cudaMemset ( dev_histo, 0,
        256 * sizeof( int ) );

    // kernel launch - 2x the number of mps gave best timing
    cudaDeviceProp prop;
    cudaGetDeviceProperties ( &prop, 0 );
    int blocks = prop.multiProcessorCount;
    printf(" multiprocesseur : %d", blocks);
    histo_kernel <<<blocks*2,256>>>( dev_buffer, SIZE, dev_histo
);

    unsigned int histo[256];
    cudaMemcpy ( histo, dev_histo,
        256 * sizeof( int ),
        cudaMemcpyDeviceToHost );
}
```

```
// get stop time, and display the timing results
    cudaEventRecord ( stop, 0 );
    cudaEventSynchronize ( stop );
    float elapsedTime;
    cudaEventElapsedTime ( &elapsedTime,
        start, stop );
    printf( "Time to generate: %3.1f ms\n" , elapsedTime );

    long histoCount = 0;
    for (int i=0; i<256; i++) {
        histoCount += histo[i];
    }
    printf( "Histogram Sum: %ld\n" , histoCount );

    // verify that we have the same counts via CPU
    for (int i=0; i<SIZE; i++)
        histo[buffer[i]]--;
    for (int i=0; i<256; i++) {
        if (histo[i] != 0)
            printf ( "Failure at %d! Off by %d\n" , i, histo[i] );
    }

    cudaEventDestroy ( start );
    cudaEventDestroy ( stop );
    cudaFree ( dev_histo );
    cudaFree ( dev_buffer );
    free ( buffer );
    system("PAUSE");
    return ;
}
```


Atomics - shared memory

- Resolución del problema anterior utilizando la shared memory.
- 256 threads/blocks
 - un histograma por blocks

```
__global__ void histo_kernel( unsigned char *buffer, long size, unsigned int *histo ) {  
  
    __shared__ unsigned int temp[256];  
    temp[threadIdx.x] = 0;  
    __syncthreads();  
  
    . . . . .
```



Atomics - shared memory

- 256 threads en competición
 - obtenemos en temp un histograma por block
 - `atomicAdd(&temp[buffer[i]], 1);`

```
__global__ void histo_kernel( unsigned char *buffer, long size, unsigned int *histo ) {  
  
    __shared__ unsigned int temp[256];  
    temp[threadIdx.x] = 0;  
    __syncthreads();  
  
    // calculate the starting index and the offset to the next  
    // block that each thread will be processing  
    int i = threadIdx.x + blockIdx.x * blockDim.x;  
    int stride = blockDim.x * gridDim.x;  
    while (i < size) {  
        atomicAdd( &temp[buffer[i]], 1 );  
        i += stride;  
    }  
    . . . . .  
}
```



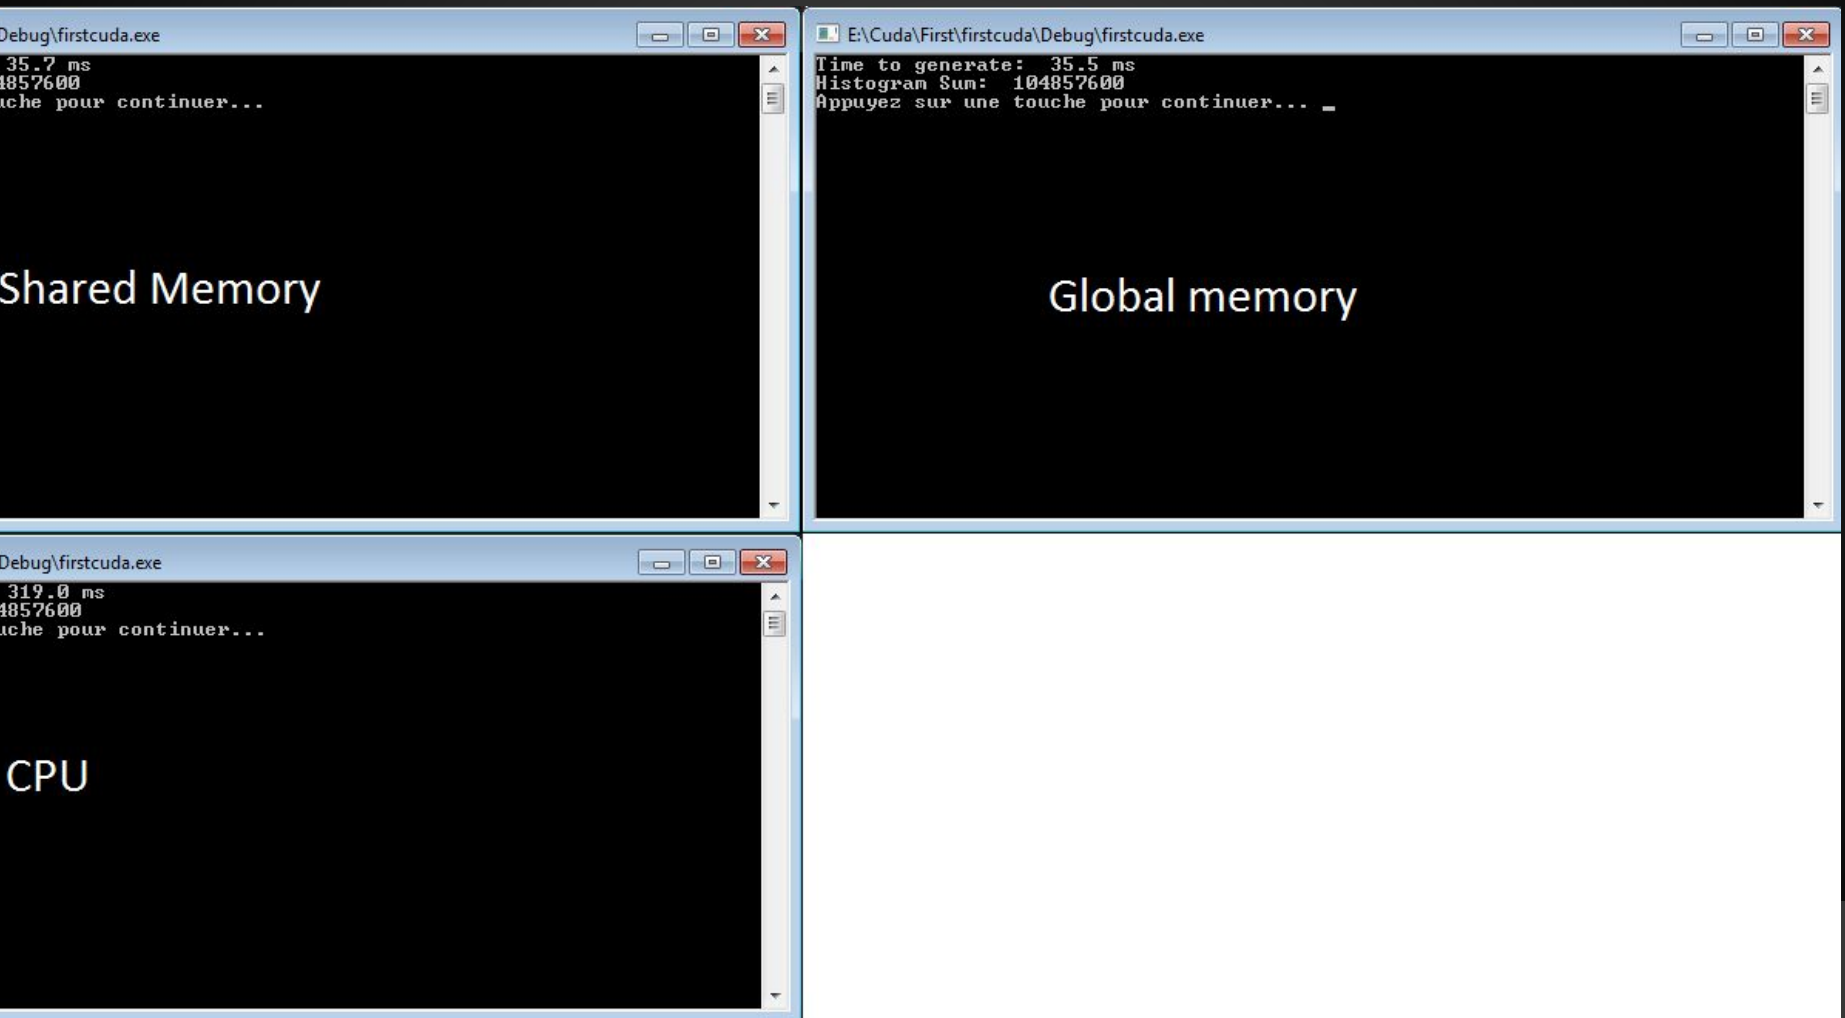
Atomics - shared memory

- cada thread tiene que sumar su resultado en histo global

- `atomicAdd(&(histo[threadIdx.x]), temp[threadIdx.x]);`

```
__global__ void histo_kernel( unsigned char *buffer, long size, unsigned int *histo ) {  
  
    __shared__ unsigned int temp[256];  
    temp[threadIdx.x] = 0;  
    __syncthreads();  
  
    // calculate the starting index and the offset to the next  
    // block that each thread will be processing  
    int i = threadIdx.x + blockIdx.x * blockDim.x;  
    int stride = blockDim.x * gridDim.x;  
    while (i < size) {  
        atomicAdd( &temp[buffer[i]], 1 );  
        i += stride;  
    }  
    //updating the  
    // global histogram is just one write per thread!  
    __syncthreads();  
    atomicAdd( &(histo[threadIdx.x]), temp[threadIdx.x] );  
}
```

Atoms - Ejemplo : Histograma



Atomics

- Resuelven problemas de competición entre thread para operación de read-modify-write
- Tener cuidado al rendimiento del programa ejecutado en GPU ya que podemos enfrentarnos con problemas de serialización de operaciones

. functions

- `atomicAdd()`
- `atomicSub()`
- `atomicMin()`
- `atomicMax()`
- ...

<http://docs.nvidia.com/cuda/cuda-c-programming-guide/index.html#arithmetic-functions>

